

Customer Segmentation Analysis

Nicholas Chavez

June 6th, 2025

Business Task: Using Online Retail data donated to UC Irvine on 11/5/2015 that recorded all transactions occurring between 01/12/2010 to 19/12/2011 for a UK-based registered non-store online retailer, I am prompted with the task of differentiating customer types to drive targeted marketing campaigns.

Setting up packages and loading data

```
library(tidyverse)
library(lubridate)
library(plotly)
setwd("C:/Users/Nicho/OneDrive/Desktop/Customer Segmentation")
sales_data <- read.csv("Online_Retail.csv")
```

Explore Data

```
head(sales_data)
```

	InvoiceNo	StockCode	Description	Quantity
1	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6
2	536365	71053	WHITE METAL LANTERN	6
3	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8
4	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6
5	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6
6	536365	22752	SET 7 BABUSHKA NESTING BOXES	2

	InvoiceDate	UnitPrice	CustomerID	Country
--	-------------	-----------	------------	---------

```

1 12/1/2010 8:26      2.55      17850 United Kingdom
2 12/1/2010 8:26      3.39      17850 United Kingdom
3 12/1/2010 8:26      2.75      17850 United Kingdom
4 12/1/2010 8:26      3.39      17850 United Kingdom
5 12/1/2010 8:26      3.39      17850 United Kingdom
6 12/1/2010 8:26      7.65      17850 United Kingdom

```

```
tibble(arrange(sales_data, Quantity))
```

```

# A tibble: 541,909 x 8
  InvoiceNo StockCode Description      Quantity InvoiceDate UnitPrice CustomerID
  <chr>     <chr>   <chr>          <int> <chr>        <dbl>       <int>
1 C581484  23843   "PAPER CRAFT ,~ -80995 12/9/2011 ~    2.08       16446
2 C541433  23166   "MEDIUM CERAMI~ -74215 1/18/2011 ~    1.04       12346
3 556690   23005   "printing smud~ -9600   6/14/2011 ~    0          NA
4 556691   23005   "printing smud~ -9600   6/14/2011 ~    0          NA
5 C536757  84347   "ROTATING SILV~ -9360   12/2/2010 ~    0.03       15838
6 556687   23003   "Printing smud~ -9058   6/14/2011 ~    0          NA
7 546152   72140F  "throw away"  -5368   3/9/2011 1~    0          NA
8 573596   79323W  "Unsaleable, d~ -4830   10/31/2011~    0          NA
9 566768   16045   ""           -3667   9/14/2011 ~    0          NA
10 565304  16259   ""           -3167   9/2/2011 1~    0          NA
# i 541,899 more rows
# i 1 more variable: Country <chr>

```

Customer ID is a unique identifier that can be used to create a summary data for frequency, InvoiceDate can be used to determine how recent orders are, and the sum of the product of Quantity and Unit price per CustomerID can be used to determine monetary spending habits per customer. Also, Invoice data is in character format, using lubridate is necessary to convert the data to date format. Additionally, Quantity can be negative, based on description of data seems to be salvage/unsaleable merchandise, this will have to be removed in the cleaning process.

Data Cleaning and mutation:

In this cleaning, I dropped null values for CustomerIDs, removed negative values for Quantity, added the TotalPurchase column which is the sum of the products of Quantity and UnitPrice, and reformatted InvoiceDate from Character String to Date.

```

cleaned_sales <- sales_data |>
  drop_na(CustomerID) |>
  filter(Quantity > 0) |>
  mutate(TotalPurchase = Quantity*UnitPrice) |>
  mutate(InvoiceDate = mdy_hm(InvoiceDate))

```

Data for Recency, Frequency, and Monetary Spending

In this section, I calculated new fields Recency, Frequency, and Monetary to help determine customer type. The first step was creating a reference date to compare dates to determine recency before using n_distinct() and sum() to create frequency and monetary.

```

reference_date <- max(cleaned_sales$InvoiceDate) + days(1)
RFM_data <- cleaned_sales |>
  group_by(CustomerID) |>
  summarise(
    recency = as.numeric(reference_date - max(InvoiceDate), units = "days"),
    frequency = n_distinct(InvoiceNo),
    monetary = sum(TotalPurchase)
  )

```

Prepare for K-means

Standardize the data for comparability between different fields to prepare for K-means.

```

Standardize <- function(x){
  (x - min(x))/(max(x)-min(x))
} # Creating a function to standardize data

# Standardize all values for comparability
rfm_stz <- RFM_data |>
  select(recency, frequency, monetary) |>
  mutate(recency = Standardize(recency)) |>
  mutate(frequency = Standardize(frequency)) |>
  mutate(monetary = Standardize(monetary))

```

check results

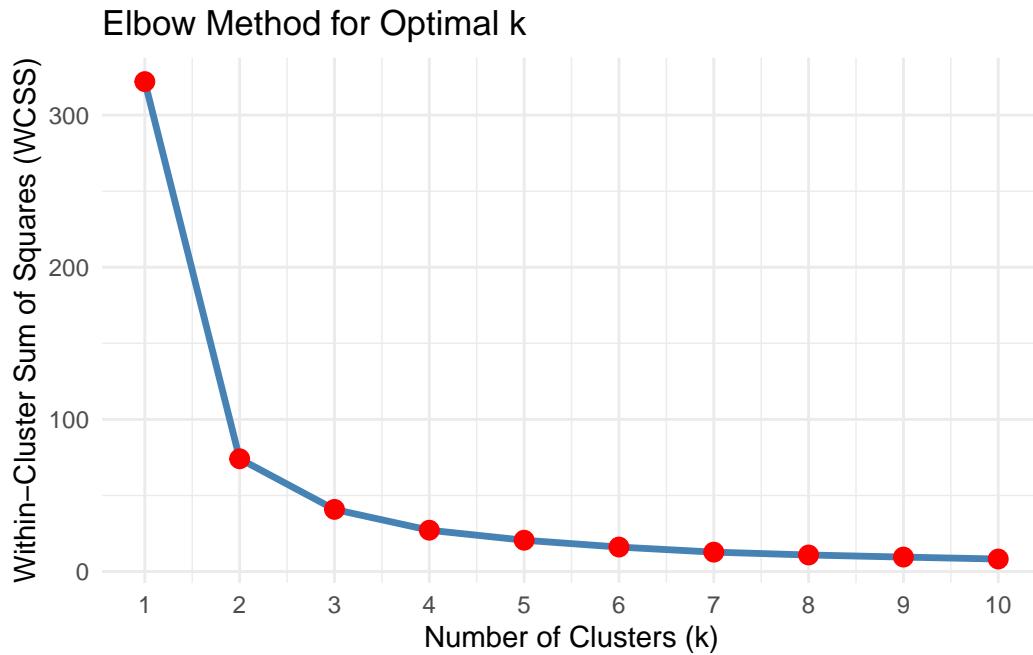
```
summary(rfm_stz)
```

	recency	frequency	monetary
Min.	:0.00000	Min. :0.000000	Min. :0.000000
1st Qu.	:0.04575	1st Qu.:0.000000	1st Qu.:0.001097
Median	:0.13423	Median :0.004785	Median :0.002407
Mean	:0.24665	Mean :0.015655	Mean :0.007330
3rd Qu.	:0.37929	3rd Qu.:0.019139	3rd Qu.:0.005930
Max.	:1.00000	Max. :1.000000	Max. :1.000000

I need to determine which the numbers of clusters to specify. For this I will be using total-within cluster sum of squares (WCSS), creating a tibble of the results, and then plot the WCSS. I then determine which k to proceed with based on the Elbow Method.

```
set.seed(123)
wcss <- map_dbl(1:10, ~{
  kmeans(rfm_stz[, c("recency", "frequency", "monetary")], centers = .,
         nstart = 25)$tot.withinss
})
wcss_data <- tibble(
  k = 1:10,
  wcss = wcss
)

ggplot(wcss_data, aes(k, wcss)) +
  geom_line(color = "steelblue", linewidth = 1.2) +
  geom_point(color = "red", size = 3) +
  labs(
    title = "Elbow Method for Optimal k",
    x = "Number of Clusters (k)",
    y = "Within-Cluster Sum of Squares (WCSS)"
  ) +
  scale_x_continuous(breaks = 1:10) +
  theme_minimal()
```



Based on the Graph, 4 is where the elbow appears and appears flat.

K-Mean Clustering

Now I will run the K-means algorithm to group/segment the customers.

```
set.seed(123)
km_out <- kmeans(rfm_stz, 4, nstart = 42)

km_out$centers
```

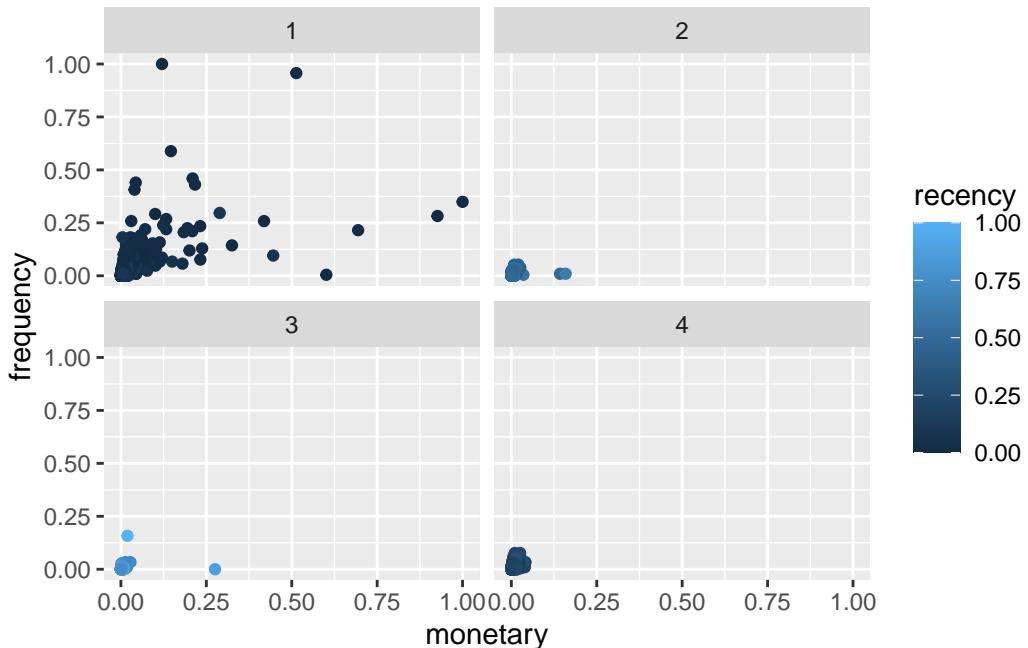
	recency	frequency	monetary
1	0.0500849	0.026288833	0.011859202
2	0.5119836	0.004180796	0.002595808
3	0.8264392	0.001671759	0.002019332
4	0.2155235	0.007792083	0.003598462

Adding the clusters to the data set

```
clustered_sales <- rfm_stz |>
  mutate(cluster = factor(km_out$cluster))
```

Visualize the data

```
ggplot(clustered_sales, aes(x = monetary, y = frequency)) +
  geom_point(aes(color = recency)) +
  facet_wrap(~cluster)
```



Findings: Through K-mean clustering I was able to determine 4 separate groups of customers. The characteristics of each group is as follows:

- **Group 1:** Mostly recent with higher spread in spending and frequency.
- **Group 2:** Third most recent with lower levels of spending and frequency. Has a few higher spenders.
- **Group 3:** Least recent with lower levels of spending and frequency.
- **Group 4:** Second most recent with lower levels of spending and frequency.

Restructure Data to prepare for further visualization

Here I will be taking the RFM_data that contains customerIDs then include the newly determine clusters before renaming them for interpretability. I will define each group based on their attributes. Group 1 will be label High Value, Group 2 would be Churning, Group 3 will be Inactive, and Group 4 are Occasional. All are defined based on the characteristics above.

```
cleaned_clusters_sales <- RFM_data |>
  mutate(Cluster = clustered_sales$cluster) |>
  mutate(
    segment = case_when(
      Cluster == 1 ~ "High Value",
      Cluster == 2 ~ "Churning",
      Cluster == 3 ~ "Inactive",
      Cluster == 4 ~ "Occasional"
    )
  )
head(cleaned_clusters_sales)
```

```
# A tibble: 6 x 6
  CustomerID recency frequency monetary Cluster segment
  <int>     <dbl>      <int>     <dbl> <fct>   <chr>
1     12346    326.        1    77184.  3     Inactive
2     12347    2.87       7     4310   1     High Value
3     12348    76.0       4     1797.  4     Occasional
4     12349    19.1       1     1758.  1     High Value
5     12350    311.       1     334.   3     Inactive
6     12352    36.9       8     2506.  1     High Value
```

Verification

I was able to apply the clusters back into the dataset because the order was the same, however that method can be prone to error. So to verify I unstandardized the values to verify they were correct. This method of verification was able to work because the ID column was drop when standardizing the data, so reverting the calculation should give the data set with the clusters the same value as the dataset with the ID.

```
unstandardize <- function(standardized_value, x){
  standardized_value*(max(x)-min(x)) + min(x)
}
check <- clustered_sales |>
```

```

  mutate(
    recency = unstandardize(recency, RFM_data$recency),
    frequency = unstandardize(frequency, RFM_data$frequency),
    monetary = unstandardize(monetary, RFM_data$monetary)
  )
head(cleaned_clusters_sales)

```

```

# A tibble: 6 x 6
  CustomerID recency frequency monetary Cluster segment
    <int>     <dbl>      <int>     <dbl> <fct>   <chr>
1     12346    326.        1    77184.  3   Inactive
2     12347     2.87       7     4310   1   High Value
3     12348     76.0       4    1797.  4   Occasional
4     12349     19.1       1    1758.  1   High Value
5     12350     311.       1     334.  3   Inactive
6     12352     36.9       8    2506.  1   High Value

```

```
head(check)
```

```

# A tibble: 6 x 4
  recency frequency monetary cluster
    <dbl>      <dbl>      <dbl> <fct>
1 326.        1    77184.  3
2 2.87        7     4310   1
3 76.0        4    1797.  4
4 19.1        1    1758.  1
5 311.        1     334.  3
6 36.9        8    2506.  1

```

```
tail(cleaned_clusters_sales)
```

```

# A tibble: 6 x 6
  CustomerID recency frequency monetary Cluster segment
    <int>     <dbl>      <int>     <dbl> <fct>   <chr>
1     18278    74.0       1    174.  4   Occasional
2     18280    278.       1    181.  3   Inactive
3     18281    181.       1     80.8  2   Churning
4     18282     8.05      2    178.  1   High Value
5     18283     4.03      16   2095. 1   High Value
6     18287    43.1       3    1837. 1   High Value

```

```

tail(check)

# A tibble: 6 x 4
  recency frequency monetary cluster
    <dbl>      <dbl>     <dbl> <fct>
1    74.0        1     174.  4
2   278.        1     181.  3
3   181.        1     80.8  2
4    8.05       2     178.  1
5    4.03       16    2095. 1
6   43.1        3     1837. 1

```

When looking at the table I looked to see if the recency,frequency, monetary, and cluster values matched. So this checks out and the data is associate with the correct customerID's. Since the data is now confirmed to be credible I can export the data to be used in other visualization tools.

```

write.csv(cleaned_clusters_sales, "Segmented_data_Irvine.csv",
          row.names = FALSE)

```

My choice to create an CSV file before closing the analysis is because the current dataset “Segmented_data_Irvine.csv” can be used for Target Marketing. Utilizing the CustomerID section and Segement will allow the marketing team to isolate customers into their respective groups and contact them with promotions and rewards.

Customer Lifetime values

CLV = frequency * lifespan * average_spending

frequency = n_distinct(InvoiceNo) lifespan = max(InvoiceDate) - min(InvoiceDate) average_spending = mean(TotalPurchases)

```

CLV_data <- cleaned_sales |>
  group_by(CustomerID) |>
  summarize(
    frequency = n_distinct(InvoiceNo),
    lifespan = as.numeric(max(InvoiceDate) - min(InvoiceDate), units = "days"),
    avg_spending = mean(TotalPurchase)) |>
  mutate(CLV = frequency * lifespan * avg_spending) |>
  left_join(cleaned_clusters_sales |> select("CustomerID", "segment"), by = "CustomerID") |>
  head(CLV_data)

```

```
# A tibble: 6 x 6
  CustomerID frequency lifespan avg_spending      CLV segment
  <int>       <int>    <dbl>        <dbl>    <dbl> <chr>
1     12346         1        0      77184.      0 Inactive
2     12347         7       365.      23.7  60512. High Value
3     12348         4       283.      58.0  65571. Occasional
4     12349         1        0      24.1      0 High Value
5     12350         1        0      19.7      0 Inactive
6     12352         8       260.      29.5  61345. High Value
```

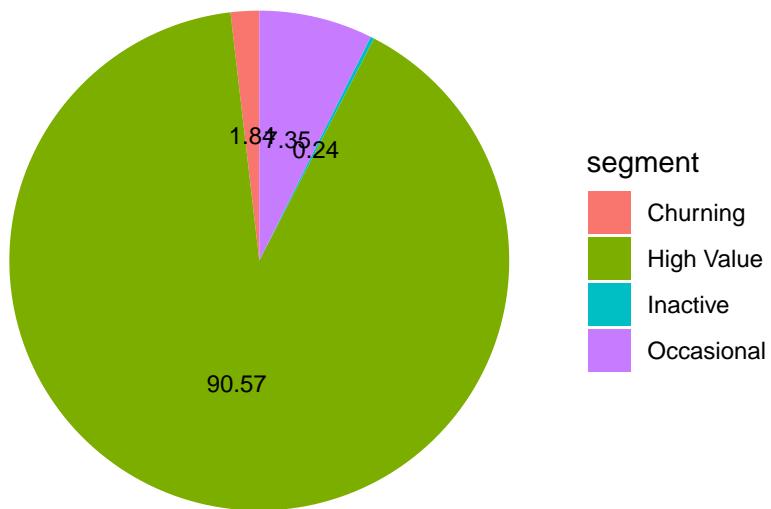
```
grouped_CLV <- CLV_data |>
  mutate( CLV_portion = (CLV / sum(CLV)) * 100 ) |>
  group_by(segment) |>
  summarize(
    avg_frequency = mean(frequency),
    avg_spending = mean(avg_spending),
    percent_of_lifetimevalue = sum(CLV_portion),
    total_value = sum(CLV)
  )
```

```
head(grouped_CLV)
```

```
# A tibble: 4 x 5
  segment      avg_frequency avg_spending percent_of_lifetimevalue total_value
  <chr>           <dbl>        <dbl>            <dbl>        <dbl>
1 Churning        1.87        64.7            1.84      4266402.
2 High Value      6.49        57.6            90.6      209751372.
3 Inactive         1.35        192.            0.241      557453.
4 Occasional      2.63        35.0            7.35      17018868.
```

```
ggplot(grouped_CLV, aes(x = "", y = percent_of_lifetimevalue, fill = segment) ) +
  geom_bar(stat = "identity", width = 1) + # Create the bar
  coord_polar("y", start = 0) +          # Wrap the bar into a circle
  theme_void() +                      # Remove background grid/axes
  labs(title = "Category Breakdown") +
  geom_text(aes(label = round(percent_of_lifetimevalue,2)),
            position = position_stack(vjust = 0.5),
            color = "black",
            size = 3)
```

Category Breakdown



Though the final calculation of customer lifetime value, and the relative percentages, we can see that 90.57% of customer lifetime value come from the high value customers, 7.35% come from occasional customers, 1.84% come from Churning customers, and 0.24% come from inactive customers.